

# lab6\_solution

November 4, 2019

## 1 COMP3222/COMP6246 Machine Learning Technologies (2019/20)

### 1.1 Lab 6 - Recurrent Neural Networks (Chapter 14)

Follow each section at your own pace, you can have a look at the book or ask questions to demonstrators if you find something confusing.

## 2 1. Basic Theory

Until now, we looked into basic perceptrons, convolutional neural network (CNN) and how to implement them in TensorFlow. In practice these techniques are used in tasks such as: searching images, self-driving cars, automatic video classification and many more. Surely, there are different network architectures that are used in Deep Learning. In the previous lab, we showed that CNNs are essentially for "processing a grid of values". However, the Deep Learning community has also generated another architecture specifically for "processing a sequence of values", which are called **Recurrent Neural Networks (RNN)** [Goodfellow 2016]. In practice, recurrent neural networks are used for analyzing time series: stock prices, car trajectories, sentiment analysis and more.

*Get Motivated:* Have a look at [this interactive example](#), which generates new strokes in your handwriting style using RNNs. The model is explained in [this paper](#).

### 2.1 Bare-bones RNN

Let's implement an RNN with five recurrent neurons without using TensorFlow's RNN implementation/utilities.

```
[1]: import numpy as np
import matplotlib.pyplot as plt

# These two lines are required to use Tensorflow 1
import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()

# to make this notebook's output stable across runs
```

```

def reset_graph(seed=42):
    tf.reset_default_graph()
    tf.set_random_seed(seed)
    np.random.seed(seed)

reset_graph()
# Let's assume some artificial data with three input (if our objective is to
→predict words in a sentence
n_inputs = 3 # then for instance: first word, second word, third word can be
→the input of our model)
n_neurons = 5 # number of neurons

X0 = tf.placeholder(tf.float32, [None, n_inputs]) # t=0 batch
X1 = tf.placeholder(tf.float32, [None, n_inputs]) # t=1 batch

# Weights on inputs (all steps share this), initially they are set random
Wx = tf.Variable(tf.random_normal(shape=[n_inputs, n_neurons], dtype=tf.float32))

# Connection weights for the outputs of the previous timestep (all steps share
→this), initially they are set random
Wy = tf.Variable(tf.random_normal(shape=[n_neurons, n_neurons], dtype=tf.float32))

# bias vector, all zeros for now
b = tf.Variable(tf.zeros([1, n_neurons], dtype=tf.float32))

# outputs of timestep 0
Y0 = tf.tanh(tf.matmul(X0, Wx) + b)

# outputs of timestep 1
Y1 = tf.tanh(tf.matmul(Y0, Wy) + tf.matmul(X1, Wx) + b)
# Y1 = activation_function(dot_product(Y0, Wy) + dot_product(X1, Wx) +
→bias_vector)

init = tf.global_variables_initializer()

# Mini-batch:           instance1  instance2  instance3 instance4
X0_batch = np.array([[0, 1, 2], [3, 4, 5], [6, 7, 8], [9, 0, 1]]) # t = 0
X1_batch = np.array([[9, 8, 7], [0, 0, 0], [6, 5, 4], [3, 2, 1]]) # t = 1

# within the session
with tf.Session() as sess:
    init.run()
    # get the outputs of each step
    Y0_val, Y1_val = sess.run([Y0, Y1], feed_dict={X0: X0_batch, X1: X1_batch})

```

WARNING:tensorflow:From C:\Local\anaconda3\envs\MLTech\lib\site-packages\tensorflow\_core\python\compat\v2\_compat.py:65:

disable\_resource\_variables (from tensorflow.python.ops.variable\_scope) is deprecated and will be removed in a future version.

Instructions for updating:

non-resource variables are not supported in the long term

```
[2]: print(Y0_val) # layers output at t=0
```

```
[[ -0.0664006   0.9625767   0.68105793   0.7091854  -0.898216   ]
 [  0.9977755  -0.71978897 -0.9965761   0.9673924  -0.9998972   ]
 [  0.99999774 -0.99898803 -0.9999989   0.9967762  -0.9999999   ]
 [  1.          -1.          -1.          -0.99818915  0.9995087   ]]
```

```
[3]: print(Y1_val) # layers output at t=1
```

```
[[ 1.          -1.          -1.          0.40200275 -0.9999998   ]
 [-0.12210423  0.6280527   0.9671843  -0.9937122  -0.25839362]
 [ 0.9999983  -0.9999994  -0.9999975  -0.8594331  -0.9999881   ]
 [ 0.99928284 -0.99999803 -0.9999058   0.9857963  -0.92205757]]
```

For the given example above, from the comments in the code:

**Exercise 1.2.: How would you define the outputs?**

Have a look at page 386 of your book, Figure 14\_1. These are the outputs of each neuron in an RNN.

**Exercise 1.3.: Why are there five columns?**

Because, there are five neurons.

**Exercise 1.4.: Why are there four rows?**

You are providing four training samples of the data.

**Exercise 1.5.: What would be the difference between `instance1` at  $t = 0$  and `instance1` at  $t = 1$ ?**

They are not related. In this example, they are two batches of different data points.

**Exercise 1.6.: What is the difference between `instance1` and `instance2` at  $t = 1$ ?**

We are supplying four training samples in each two batches. Therefore, they are just different data points.

## 3 2. Predicting Time Series

Let's look at a simple use of RNNs with time series, these time series could be stock prices, brain wave patterns and so on. Our objective could be predicting the future stock price, given the available data that we have.

Let's define an arbitrary sine function for stock prices `time_series(t)` to make our predictions.

```

[4]: reset_graph()
      # time starts from 0 to 30
      t_min, t_max = 0, 30
      # we sample time_series function for every 0.1
      resolution = 0.1

      def time_series(t):
          return t * np.sin(t) / 3 + 2 * np.sin(t*5)

      def next_batch(batch_size, n_steps):
          """
          Returns a batch with `n_steps`: number of instances
          """
          # randomly get a starting number between a range
          t0 = np.random.rand(batch_size, 1) * (t_max - t_min - n_steps * resolution)
          # make a list until of number with n_steps until the next batch
          Ts = t0 + np.arange(0., n_steps + 1) * resolution
          # get the outputs of time_series function given the input Ts (time points)
          ys = time_series(Ts)

          # return X's and Y's
          return ys[:, :-1].reshape(-1, n_steps, 1), ys[:, 1:].reshape(-1, n_steps, 1)

      # inputs to the time_series function
      t = np.linspace(t_min, t_max, int((t_max - t_min) / resolution))

      n_steps = 20
      # a training instance
      t_instance = np.linspace(12.2, 12.2 + resolution * (n_steps + 1), n_steps + 1)

      plt.figure(figsize=(11,4))
      plt.subplot(121)
      plt.title("A time series (generated)", fontsize=14)
      # plot all the data
      plt.plot(t, time_series(t), label=r"$t \cdot \sin(t) / 3 + 2 \cdot \sin(5t)$")

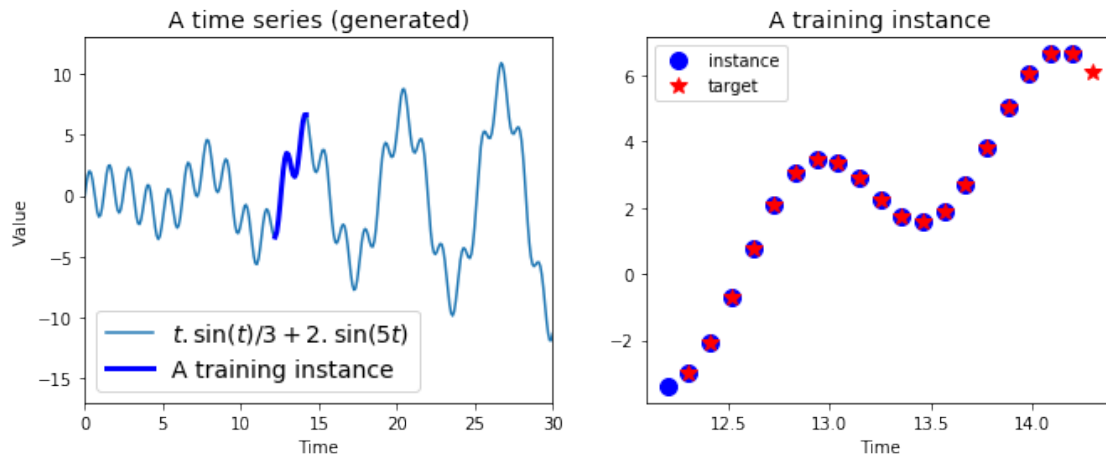
      # plot only the training set
      plt.plot(t_instance[:-1], time_series(t_instance[:-1]), "b-", linewidth=3,
               label="A training instance")
      plt.legend(loc="lower left", fontsize=14)
      plt.axis([0, 30, -17, 13])
      plt.xlabel("Time")
      plt.ylabel("Value")

      plt.subplot(122)
      plt.title("A training instance", fontsize=14)

```

```
plt.plot(t_instance[:-1], time_series(t_instance[:-1]), "bo", markersize=10,
        ↪label="instance")
# notice that targets are shifted by one time step into the future
plt.plot(t_instance[1:], time_series(t_instance[1:]), "r*", markersize=10,
        ↪label="target")
plt.legend(loc="upper left")
plt.xlabel("Time")

plt.show()
```



```
[5]: X_batch, y_batch = next_batch(1, n_steps)

# combining X_batch and y_batch for better printing, first_column=X,
↪second_column=Y
print(np.c_[X_batch[0], y_batch[0]])
# Did you notice the shift in y values?
```

```
[[-1.40208096 -2.33035999]
 [-2.33035999 -3.4513234 ]
 [-3.4513234  -4.52641909]
 [-4.52641909 -5.32081479]
 [-5.32081479 -5.66045846]
 [-5.66045846 -5.47433377]
 [-5.47433377 -4.81157012]
 [-4.81157012 -3.82922233]
 [-3.82922233 -2.75371563]
 [-2.75371563 -1.82539786]
 [-1.82539786 -1.23977629]
 [-1.23977629 -1.0998269 ]
 [-1.0998269  -1.39105208]
 [-1.39105208 -1.98539218]]
```

```

[-1.98539218 -2.67303091]
[-2.67303091 -3.214304 ]
[-3.214304   -3.39899794]
[-3.39899794 -3.09851497]
[-3.09851497 -2.29812628]
[-2.29812628 -1.10140997]]

```

```

[6]: reset_graph()

n_steps = 20
n_inputs = 1
n_neurons = 100
n_outputs = 1
learning_rate = 0.001
n_iterations = 1500
batch_size = 50

# Optimizer that finds the weight values for each neuron
def get_predictions(optimizer="gdo",
                    loss_function="mse",
                    save=False,
                    reset=True):

    reset_graph()
    X = tf.placeholder(tf.float32, [None, n_steps, n_inputs])
    y = tf.placeholder(tf.float32, [None, n_steps, n_outputs])

    # We use `dynamic_rnn` and `BasicRNNCell` utilities in this case, with tf.
    →nn.relu
    cell = tf.nn.rnn_cell.BasicRNNCell(num_units=n_neurons, activation=tf.nn.
    →relu)
    rnn_outputs, states = tf.nn.dynamic_rnn(cell, X, dtype=tf.float32)

    # This part is visually shown in the book Figure 14-10.
    stacked_rnn_outputs = tf.reshape(rnn_outputs, [-1, n_neurons])

    # What do you think line below will be doing? (Tip: https://www.tensorflow.
    →org/api_docs/python/tf/layers/dense)
    stacked_outputs = tf.layers.dense(stacked_rnn_outputs, n_outputs)

    outputs = tf.reshape(stacked_outputs, [-1, n_steps, n_outputs])

    if optimizer == "gdo":
        optimizer = tf.train.
    →GradientDescentOptimizer(learning_rate=learning_rate)
    elif optimizer == "adam":
        optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)

```

```

# The loss function to optimize
if loss_function == "mse":
    loss = tf.reduce_mean(tf.square(outputs - y))
elif loss_function == "rmse":
    loss = tf.sqrt(tf.reduce_mean(tf.squared_difference(outputs, y)))

# Let the optimizer know that this is the loss function to optimize
training_op = optimizer.minimize(loss)

init = tf.global_variables_initializer()
saver = tf.train.Saver()
y_pred = None
with tf.Session() as sess:
    init.run()
    for iteration in range(n_iterations):
        # get a random batch
        X_batch, y_batch = next_batch(batch_size, n_steps)
        # run tensorflow session
        sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
        # in each 100th iteration
        if iteration % 100 == 0: # with RSME
            loss_val = loss.eval(feed_dict={X: X_batch, y: y_batch})
            # print the MSE
            print(iteration, "\t{}:".format(loss_function), loss_val)

    X_new = time_series(np.array(t_instance[:-1].reshape(-1, n_steps,
↪n_inputs)))
    y_pred = sess.run(outputs, feed_dict={X: X_new})
    if save:
        saver.save(sess, "./my_time_series_model_" + loss_function)
    if reset:
        reset_graph()

return y_pred, saver, outputs, X

```

```

[7]: y_mse_pred, _, _, _ = get_predictions(optimizer="gdo", loss_function="mse")
y_rmse_pred, _, _, _ = get_predictions(optimizer="gdo", loss_function="rmse",
↪save=True)
y_mse_pred, y_rmse_pred

```

WARNING:tensorflow:From <ipython-input-6-ffd5b8170bb1>:22: BasicRNNCell.\_\_init\_\_ (from tensorflow.python.ops.rnn\_cell\_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This class is equivalent as tf.keras.layers.SimplerNNCell, and will be replaced

by that in Tensorflow 2.0.

WARNING:tensorflow:From <ipython-input-6-ffd5b8170bb1>:23: dynamic\_rnn (from tensorflow.python.ops.rnn) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `keras.layers.RNN(cell)`, which is equivalent to this API

WARNING:tensorflow:From C:\Local\anaconda3\envs\MLTech\lib\site-packages\tensorflow\_core\python\ops\rnn\_cell\_impl.py:456: Layer.add\_variable (from tensorflow.python.keras.engine.base\_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `layer.add\_weight` method instead.

WARNING:tensorflow:From C:\Local\anaconda3\envs\MLTech\lib\site-packages\tensorflow\_core\python\ops\rnn\_cell\_impl.py:460: calling Zeros.\_\_init\_\_ (from tensorflow.python.ops.init\_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From <ipython-input-6-ffd5b8170bb1>:29: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.

Instructions for updating:

Use keras.layers.Dense instead.

WARNING:tensorflow:From C:\Local\anaconda3\envs\MLTech\lib\site-packages\tensorflow\_core\python\layers\core.py:187: Layer.apply (from tensorflow.python.keras.engine.base\_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `layer.\_\_call\_\_` method instead.

0	mse: 13.841026
100	mse: 0.9208843
200	mse: 0.5451863
300	mse: 0.36086816
400	mse: 0.29775804
500	mse: 0.26228595
600	mse: 0.23783027
700	mse: 0.1819096
800	mse: 0.17917499
900	mse: 0.18779245
1000	mse: 0.16793889
1100	mse: 0.16937979
1200	mse: 0.13638055
1300	mse: 0.14248464
1400	mse: 0.116242446
0	rmse: 3.8650882
100	rmse: 1.8601373
200	rmse: 1.1923529



```

300    rmse: 0.95981115
400    rmse: 0.7683989
500    rmse: 0.6816613
600    rmse: 0.61443514
700    rmse: 0.5278193
800    rmse: 0.5077234
900    rmse: 0.5024839
1000   rmse: 0.46223927
1100   rmse: 0.45384395
1200   rmse: 0.40254924
1300   rmse: 0.401823
1400   rmse: 0.36164722

```

```

[7]: (array([[ -3.5477045 ],
             [ -3.0211253 ],
             [ -1.3003355 ],
             [  0.50970757],
             [  2.3014107 ],
             [  3.040976  ],
             [  3.4921103 ],
             [  3.4513597 ],
             [  2.7930527 ],
             [  2.229236  ],
             [  1.54554   ],
             [  1.2409023 ],
             [  1.5544744 ],
             [  2.3978558 ],
             [  3.6025336 ],
             [  4.716492  ],
             [  5.714428  ],
             [  6.556115  ],
             [  6.739255  ],
             [  6.2484055 ]]), dtype=float32), array([[ -3.508388 ],
             [ -3.0818055],
             [ -1.3413746],
             [  0.5104261],
             [  2.268932  ],
             [  3.0400412],
             [  3.4817607],
             [  3.4555745],
             [  2.7951434],
             [  2.2374256],
             [  1.5505079],
             [  1.2343997],
             [  1.5314096],
             [  2.3610237],
             [  3.5536125],


```


```
[ 4.671324 ],
[ 5.7337027],
[ 6.549388 ],
[ 6.7378163],
[ 6.26523  ]]], dtype=float32))
```

```
[8]: def plot_results(title, y_vals):
    plt.title(title, fontsize=14)
    plt.plot(t_instance[:-1], time_series(t_instance[:-1]), "bo",
    ↪markersize=10, label="instance")
    plt.plot(t_instance[1:], time_series(t_instance[1:]), "r*", markersize=10,
    ↪label="target")
    plt.plot(t_instance[1:], y_vals[0,:,0], "c.", markersize=10,
    ↪label="prediction")
    plt.legend(loc="upper left")
    plt.xlabel("Time")

    plt.show()

plot_results("mse", y_mse_pred)
plot_results("rmse", y_rmse_pred)
```





**Exercise 2.1.** Add comments to the code blocks above. Do you understand the purpose of each line?

Check the code blocks above.

**Exercise 2.2.** How can you improve the MSE? (*Tip: Remember Lab 4: Gradient Descent*)

The code blocks above are updated with a lower learning rate and higher epoch time.

**Exercise 2.3.** Implement the RMSE instead of the MSE, compare the test plots.

Implemented above. Difference is quite minimal in the plots.

## 4 3. Generative RNNs

We can use RNNs to generate sequences, below you are going to use the model we trained. You should expect some resemblance to the original time series.

```
[9]: # since the graph is reset, let's train again:
loss_function = "mse"
_, saver, outputs, X = get_predictions(optimizer="adam",
    ↪ loss_function=loss_function, save=True, reset=False)

with tf.Session() as sess:
```

```

saver.restore(sess, "./my_time_series_model_"+loss_function)

sequence = [0.] * n_steps
for iteration in range(300):
    X_batch = np.array(sequence[-n_steps:]).reshape(1, n_steps, 1)
    y_pred = sess.run(outputs, feed_dict={X: X_batch})
    sequence.append(y_pred[0, -1, 0])

plt.figure(figsize=(11,4))
plt.subplot(121)
plt.title("A time series (generated)", fontsize=14)
# plot all the data
plt.plot(t, time_series(t), label=r"$t \cdot \sin(t) / 3 + 2 \cdot \sin(5t)$")

plt.legend(loc="lower left", fontsize=14)
plt.axis([0, 30, -17, 13])
plt.xlabel("Time")
plt.ylabel("Value")

plt.subplot(122)
plt.title("Time series generated by RNN", fontsize=14)

plt.plot(np.arange(len(sequence)), sequence, "b-")
plt.plot(t[:n_steps], sequence[:n_steps], "b-", linewidth=3)
plt.xlabel("Time")
plt.ylabel("Value")


plt.show()

```

```

0      mse: 13.907031
100    mse: 0.50563276
200    mse: 0.1972298
300    mse: 0.1026313
400    mse: 0.067478806
500    mse: 0.06297704
600    mse: 0.05659878
700    mse: 0.050051264
800    mse: 0.0505437
900    mse: 0.04874706
1000   mse: 0.04847027
1100   mse: 0.050347283
1200   mse: 0.041849542
1300   mse: 0.05084179
1400   mse: 0.043897416
INFO:tensorflow:Restoring parameters from ./my_time_series_model_mse

```



**Exercise 3.1.** Does your plot resemble the actual time series? Why do you think so?

You should notice that if you improve your errors, the time series generated would resemble the actual time series.

**Exercise 3.2.** Change your optimizer to `AdamOptimizer`, what do you think has changed?

The code is now more structured, please try with different parameters.

**Exercise 3.3.** Try different activation functions. (e.g. `logit`, `tanh`, ...)

Have a look at the `generate_predictions` function. Follow the same pattern and implement the other activation function and compare the time series & errors.

## 5 Recap

In this lab, we demonstrated these concepts:

- from theory to implementation, how a simple RNN works
- how to predict a time series with RNN
- which parameters to look out for in order to improve the predictions
- generation of sequences with a RNN

As in the previous labs, there is some material that we have not been able to cover. In your free time, you can have a look at:

- LSTM Cells and GRU Cells
- NLP Applications with RNNs
- Encoding and Decoding with RNNs

### 5.0.1 References

[Goodfellow, 2016] : <https://www.deeplearningbook.org/>